

Applications of LLM – Part I: Retrieval Augmented Generation (RAG)

Week 9 - LGT

Dr Lin Gui

Lin.1.gui@kcl.ac.uk



Learning outcomes


- By the end of this topic, you will be able to:
 - Understand the core concepts of Retrieval-Augmented Generation and how it differs from standard LLM approaches.
 - Build and configure a basic RAG pipeline using embeddings, retrievers, and generators.
 - Evaluate and optimize RAG performance through effective data preparation, chunking, and retrieval strategies.

Contents

- RAG overview
- Foundation of information Retrieval
- RAG Paradigms Shifting
- Key Technologies and Evaluation
- Applications


RAG overview

- When answering questions or generating text, it first retrieves relevant information from a large number of documents, and then LLMs generates answers based on this information.
- By attaching an external knowledge base, there is no need to retrain the entire large model for each specific task.
- The RAG model is especially suitable for knowledge-intensive tasks.



Symbolic Knowledge or Parametric Knowledge

- Ways to optimize LLMs.
 - Prompt Engineering
 - Instruct / Fine-tuning
 - Retrieval-Augmented Generation



RAG vs Fine-tuning

Feature Comparison	RAG	Fine-Tuning
Knowledge Updates	Directly updating the retrieval knowledge base ensures that the information remains current without the need for frequent retraining, making it well-suited for dynamic data environments.	Stores static data, requiring retraining for knowledge and data updates.
External Knowledge	Proficient in leveraging external resources, particularly suitable for accessing documents or other structured/unstructured databases.	Can be utilized to align the externally acquired knowledge from pretraining with large language models, but may be less practical for frequently changing data sources.
Data Processing	Involves minimal data processing and handling.	Depends on the creation of high-quality datasets, and limited datasets may not result in significant performance improvements.
Model Customization	Focuses on information retrieval and integrating external knowledge but may not fully customize model behavior or writing style.	Allows adjustments of LLM behavior, writing style, or specific domain knowledge based on specific tones or terms.
Interpretability	Responses can be traced back to specific data sources, providing higher interpretability and traceability.	Similar to a black box, it is not always clear why the model reacts a certain way, resulting in relatively lower interpretability.
Computational Resources	Depends on computational resources to support retrieval strategies and technologies related to databases. Additionally, it requires the maintenance of external data source integration and updates.	The preparation and curation of high-quality training datasets, defining fine-tuning objectives, and providing corresponding computational resources are necessary.
Latency Requirements	Involves data retrieval, which may lead to higher latency.	LLM after fine-tuning can respond without retrieval, resulting in lower latency.
Reducing Hallucinations	Inherently less prone to hallucinations as each answer is grounded in retrieved evidence.	Can help reduce hallucinations by training the model based on specific domain data but may still exhibit hallucinations when faced with unfamiliar input.
Ethical and Privacy Issues	Ethical and privacy concerns arise from the storage and retrieval of text from external databases.	Ethical and privacy concerns may arise due to sensitive content in the training data.

RAG Application

- Scenarios where RAG is applicable:
 - Long-tail distribution of data
 - Frequent knowledge updates
 - Answers requiring verification and traceability
 - Specialized domain knowledge
 - Data privacy preservation

Question Answering

Fact checking

Dialog systems

Summarisation

Machine translation

Code generation

Sentiment Analysis

Commonsense
reasoning

Contents

- RAG overview
- Foundation of information Retrieval
- RAG Paradigms Shifting
- Key Technologies and Evaluation
- Applications

Foundation of information Retrieval

- What is information Retrieval?
 - The system searches collections for items relevant to the user's query. It then returns those items to the user, typically in list form sorted per computed relevance[#]
- Three main questions in information retrieval:
 - How to map the text into features (**Embedding method**)
 - How to measure the similarity between features (**IR Modelling**)
 - How to do it efficiently (**Indexing**)

Embedding Method

- How to map the text into features (vectors)?
 - Discrete representation
 - Convert the input query/document into vectors based on the lexicon
 - Continuous representation
 - Using representation learning to convert the input text into vectors

Discrete representation

- In discrete representation, for both query and document, we assign each word a specific dimension. If a word appears query/document, then value of the corresponding dimension is:
 - In Binary representation: 1
 - In TF (term frequency) based representation: t (how many times this word appears within the query/documents)
 - In TF-IDF (inverse document frequency) based representation: $t \log(n/x)$
 - Here, t is term frequency, n is number of documents, x is the number of documents which contains this term.

Discrete representation (example)

- We have the following documents:
 - D1 = “Shipment of gold damaged in a fire”.
 - D2 = “Delivery of silver arrived in a silver truck”.
 - D3 = “Shipment of gold arrived in a truck”
- After pre-processing:
 - D1 = “shipment”, “gold”, “damage”, “fire”.
 - D2 = “delivery”, “silver”, “arrive”, “truck”.
 - D3 = “shipment”, “gold”, “arrive”, “truck”.

Discrete representation (example)

- Building vocabulary:
 - $V = \text{"shipment", "gold", "damage", "fire", "delivery", "silver", "arrive", "truck"}$.
- Detect the feature for each document. If the feature occurs, the corresponding value is '1', otherwise '0' (binary feature):

	shipment	gold	damage	fire	delivery	silver	arrive	truck
D1	1	1	1	1	0	0	0	0
D2	0	0	0	0	1	1	1	1
D3	1	1	0	0	0	0	1	1

Discrete representation (example)

- **Definition – term frequency (TF):**
 - t - how many times the term appears in the document
- **Example:**
 - D1 = “Shipment of gold damaged in a fire”.
 - D2 = “Delivery of silver arrived in a silver truck”.
 - D3 = “Shipment of gold arrived in a truck”

	shipment	gold	damage	fire	delivery	silver	arrive	truck
D1	1	1	1	1	0	0	0	0
D2	0	0	0	0	1	2	1	1
D3	1	1	0	0	0	0	1	1

Discrete representation (example)

- **Definition – inverse document frequency (IDF):**
 - $\log(n/x)$ – n is number of documents, x is the number of documents which contains this term
- **Example:**
 - D1 = “Shipment of gold damaged in a fire”.
 - D2 = “Delivery of silver arrived in a silver truck”.
 - D3 = “Shipment of gold arrived in a truck”

shipment	gold	damage	fire	delivery	silver	arrive	truck
0.176	0.176	0.477	0.477	0.477	0.477	0.176	0.176

Inverse document frequency vector

Discrete representation (example)

	shipment	gold	damage	fire	delivery	silver	arrive	truck
D1	1	1	1	1	0	0	0	0
D2	0	0	0	0	1	2	1	1
D3	1	1	0	0	0	0	1	1

Term frequency matrix

shipment	gold	damage	fire	delivery	silver	arrive	truck
0.176	0.176	0.477	0.477	0.477	0.477	0.176	0.176

Inverse document frequency vector

	shipment	gold	damage	fire	delivery	silver	arrive	truck
D1	0.176	0.176	0.477	0.477	0	0	0	0
D2	0	0	0	0	0.477	0.954	0.176	0.176
D3	0.176	0.176	0	0	0	0	0.176	0.176

TF-IDF Matrix

Embedding Method

- How to map the text into features (vectors)?
 - Discrete representation
 - Convert the input query/document into vectors based on the lexicon
 - **Continuous representation**
 - **Using representation learning to convert the input text into vectors**

Continuous representation


- Continuous representation - using representation learning to convert the input text into vectors
 - Define the relation first, then using optimiser to update the embedding to approximate the relation.

Continuous representation

- Continuous representation - using representation learning to convert the input text into vectors
 - Define the relation first, then using optimiser to update the embedding to approximate the relation.

Dense Passage Retrieval

- Encode questions and text passages into continuous vectors (embeddings) and retrieve passages using vector similarity instead of keyword overlapping.
- Train directly on question–passage pairs, using **in-batch negatives** to improve efficiency.



In each batch, there are multiple question–answer pairs, both matched and unmatched. Matched pairs should have similar representations, while unmatched pairs should have representations that are far apart.

Training phase
Fine-tune two encoders

ReContriever

- What if we don't have annotated data (Matched and unmatched QA-pair).
- Using pseudo-examples: For each passage/document p , create an augmented version p' . Then treat (p, p') as a positive pair:
 - Masking words (random word masking)
 - Span deletion
 - Back-translation Sentence
 - Reordering Adding noise
 - Perturbations Cropping (taking a subset of sentences)

Using API

- There are many APIs could do this job, for example, Mistral AI:

The screenshot shows a section of the Mistral Embed API documentation titled "How to Generate Embeddings". It includes a brief description of how to use the API endpoint with the "mistral-embed" model to generate numerical vectors from input texts. Below the description are code snippets in Python, TypeScript, and CURL. The Python snippet is as follows:

```
import os
from mistralai import Mistral

api_key = os.environ["MISTRAL_API_KEY"]
model = "mistral-embed"

client = Mistral(api_key=api_key)

embeddings_batch_response = client.embeddings.create(
    model=model,
    inputs=["Embed this sentence.", "As well as this one."]
)
```

The output is an embedding object with the embeddings and the token usage information.

Let's take a look at the length of the first embedding:

```
PYTHON    TYPESCRIPT    CURL
len(embeddings_batch_response.data[0].embedding)
```

IR Modelling

- What is information Retrieval?
 - The system searches collections for items relevant to the user's query. It then returns those items to the user, typically in list form sorted per computed relevance
- Three main questions in information retrieval:
 - How to map the text into features (**Embedding method**)
 - **How to measure the similarity between features (IR Modelling)**
 - How to do it efficiently (**Indexing**)

IR Modelling

- In IR modelling, we use different metric to measure the similarity/distance between the query and given documents. The target is to find the top-k relevant documents based on the given query.
 - Cosine similarity (for both Discrete & Continuous representation)
 - Jaccard distance (for Discrete representation only)
 - BM25 (for Discrete representation only)

Cosine similarity

- Cosine similarity

$$Cos(x, y) = \frac{\sum_{i=1}^n x_i y_i}{\sqrt{\sum_{i=1}^n (x_i)^2} \sqrt{\sum_{i=1}^n (y_i)^2}}$$

- Considering

- D1 = [1,1,1,1,0,0,0,0]
- D3 = [1,1,0,0,0,0,1,1]


$$Cos(D1, D3) = 1/2$$

Jaccard similarity

- Only considering if there is overlapping or not. We don't care about the value.

- For example:

- $R_x = [2, 0, 3, 3]$
- $R_y = [1, 1, 0, 5]$



- Jaccard similarity: $\text{sim}(x, y) = \frac{|R_x \cap R_y|}{|R_x \cup R_y|}$

BM25

- BM25 is a lexicon based retrieval method that ranks a set of documents based on the query terms appearing in each document, regardless of their proximity within the document.
- Given a query Q , containing keywords q_1, q_2, \dots, q_n , the BM25 score of a document D is:

$$score(D, Q) = \sum_{i=1}^n IDF(q_i) \frac{f(q_i, D)(k_1 + 1)}{f(q_i, D) + k_1(1 - b + b \frac{|D|}{avgdl})}$$

hyperparameters

Length of D

Average length of all docs

BM25

- BM25 is a lexicon based retrieval method that ranks a set of documents based on the query terms appearing in each document, regardless of their proximity within the document.
- Given a query Q , containing keywords q_1, q_2, \dots, q_n , the BM25 score of a document D is:

$$score(D, Q) = \sum_{i=1}^n IDF(q_i) \frac{f(q_i, D)(k_1 + 1)}{f(q_i, D) + k_1(1 - b + b \frac{|D|}{avgdl})}$$

Length of D
Average length of all docs
hyperparameters



Maybe...a bit confusing
Can you speak in English?

BM25

- BM25 is a lexicon based retrieval method that ranks a set of documents based on the query terms appearing in each document, regardless of their proximity within the document.
- Given a query Q , containing keywords q_1, q_2, \dots, q_n , the BM25 score of a document D is:

$$score(D, Q) = \sum_{i=1}^n IDF(q_i) \frac{f(q_i, D)(k_1 + 1)}{f(q_i, D) + k_1(1 - b + b \frac{|D|}{avgdl})}$$

Length of D

Average length of all docs

hyperparameters

Relax...it is pretty simple actually

BM25

- BM25 is a lexicon based retrieval method that ranks a set of documents based on the query terms appearing in each document, regardless of their proximity within the document.
- Given a query Q , containing keywords q_1, q_2, \dots, q_n , the BM25 score of a document D is: **IDF term in Q (is it an important word?)**

$$score(D, Q) = \sum_{i=1}^n \frac{IDF(q_i) f(q_i, D)(k_1 + 1)}{f(q_i, D) + k_1(1 - b + b \frac{|D|}{avgdl})}$$

The ‘percentage’ of querying words in D

Indexing

- Next question, how to do it efficiently (**Indexing**)
- Suppose we have 1k queries, and there are 1 billion documents in knowledge based, how many times of comparison we need?
- $1k \times 1b$

It is a really huge number.

In real world scenario, it could be even larger

If there is only one important task in information retrieval, it must be “indexing”

Indexing - Discrete representation

- Inverted index
- Since the discrete representation is sparse (most dims are zero), we can build inverted index. For each word, we build a link list to store all the documents contain this word.
- For the given query, the complexity is now only related to the #unique words in the query. (**In most queries, the size is just few words**)

docID	geo-scopeID
1	Europe
2	Europe
3	France
4	England
5	Portugal
6	Quebec
7	Europe
8	Spain


Forward Index

geo-scopeID	docID
Europe	1 2 7
France	3
Portugal	5
England	4
Quebec	6
Spain	8

Inverted Index

Indexing - Continuous representation

- In continuous representation, it might be a bit complex. There is no sparse representation anymore.
 - We can use the following method to speed up the searching.
 - Vector compression – reduce the size of vectors
 - Hierarchical clustering – in each layer only search the nearest cluster
- Clustering the documents first, and then,
Only consider the nearest centroid during the searching**



Contents


- RAG overview
- Foundation of information Retrieval
- **RAG Paradigms Shifting**
- Key Technologies and Evaluation
- Applications

Naive RAG

- **Step 1 – indexing**
 - Divide the document into even chunks, each chunk being a piece of the original text.
 - Using the encoding model to generate an embedding for each chunk.
 - Store the Embedding of each block in the vector database.
- **Step 2 – Retrieval**
 - Retrieve the k most relevant documents using vector similarity search.
- **Step 3 – Generation**
 - The original query and the retrieved text are combined and input into a LLM to get the final answer


Naive RAG

- Step 1 – indexing
- Step 2 – Retrieval
- Step 3 – Generation



Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation



Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
 - Step 2 – Retrieval
 - +post-retrieval process
 - Step 3 – Generation
- Sliding windows

Fine-grained segmentation

Adding metadata

Advanced RAG


- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation

Sliding windows

Fine-grained segmentation

Adding metadata

Document



Split the doc into chunks, and ensure there is over lapping between chunks (WHY?)


Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation

Sliding windows

Fine-grained segmentation



Adding metadata



Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation

The metadata is the aspects of each chunk.
It will help both retriever and generator to
improve the performance.



Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation

Retrieve routes

Summarization

Rewriting

Confidence judgment

Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation

Retrieve routes = multiple retrieval paths that a RAG system can choose from, depending on query intent, data type, or document structure.




Instead of one flat “retrieve chunks by embeddings” step, you can:



Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation




Searching in summarisation
instead of full documents

Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
- Step 2 – Retrieval
 - +post-retrieval process
- Step 3 – Generation

Benefits:

a more explicit query
a more keyword-rich query
a more structured query
multiple diverse sub-queries



Advanced RAG

- Step 1 – indexing
 - + index optimization
 - + pre-retrieval process
 - Step 2 – Retrieval
 - +post-retrieval process
 - Step 3 – Generation
-
- ```
graph LR; Query[Query] --> Document[Document]; Document --> LLM[LLM]; LLM --> Confidence[Confidence checking]; Confidence --> Output[Output];
```
- The diagram illustrates the Advanced RAG process flow. It starts with a 'Query' (blue box) which leads to a 'Document' (green box). This document then feeds into an 'LLM' (purple box). The output from the LLM goes to 'Confidence checking' (orange box), which is enclosed in a dashed box. Finally, the process leads to 'Output' (blue box). A red box labeled 'Confidence judgment' is positioned above the 'Confidence checking' box.

Confirm the Confidence before output  
By similarity scores  
By LLM Confidence scores

# Advanced RAG

---

- Step 1 – indexing
  - + index optimization
  - + pre-retrieval process
- Step 2 – Retrieval
  - +post-retrieval process
- Step 3 – Generation

Re-order

Filter content retrieval

# Advanced RAG

- Step 1 – indexing
  - + index optimization
  - + pre-retrieval process
- Step 2 – Retrieval
  - +post-retrieval process
- Step 3 – Generation

Re-order

Filter content retrieval




LLMs is sensitive with the input order

The early input chunks has higher weights

How to organize the searched evidence for final output is important

# Advanced RAG

- Step 1 – indexing
  - + index optimization
  - + pre-retrieval process
- Step 2 – Retrieval
  - +post-retrieval process
- Step 3 – Generation



To avoid possible hallucination, filtering the irrelevant evidences.

# Modular RAG

---

- **Naïve RAG**



- **DSP**




- **Rewrite-Retrieve-Read**



- **Retrieve-then-read**



# Different RAG Paradigms



# Key problems in RAG

---

- How to retrieve
- When to retrieve
- How to use the retrieved information

# How to retrieve

---

- By using the information on different structuration levels

- Token level

It excels in handling long-tail and cross-domain issues with high computational efficiency, but it requires significant storage.

- Phrase level

The search is broad, recalling a large amount of information, but with low accuracy, high coverage but includes much redundant information.


- Chunk level

- Entity level

Richer semantic and structured information, but the retrieval efficiency is lower and is limited by the quality of KG.

# When to retrieve


- Two questions:
- When we need to retrieve information to support the QA
- How many times we need to retrieve the information
- Solution#1: Conducting once search during the reasoning process.



High efficiency, but low relevance of the retrieved documents

# When to retrieve


- Two questions:
- When we need to retrieve information to support the QA
- How many times we need to retrieve the information
- Solution#2: Adaptively conduct the search.



Balancing efficiency and information  
might not yield the optimal solution

# When to retrieve

- Two questions:
  - When we need to retrieve information to support the QA
  - How many times we need to retrieve the information
- 
- Solution#3: Retrieve once for every N tokens generated.



# Contents

---

- RAG overview
- RAG Paradigms Shifting
- **Key Technologies and Evaluation**
- Applications


# Key Technologies

---

- Data indexing optimization
- Structured Corpus
- Retrieval Source Optimization
- KG as a Retrieval Data Source
- Query Optimization
- Embedding Optimization
- Fine-tuning on RAG


# Data indexing optimization

- Chunk optimization
- Small-2-big: Embedding at sentence level expand the window during generation process.



# Data indexing optimization

- Chunk optimization
- Sliding window: sliding chunk covers the entire text, avoiding semantic ambiguity.




# Data indexing optimization

---

- Chunk optimization
- Two-stage method: Retrieve documents through summaries, then retrieve text blocks from the documents.
  - Step 1 — Search Summaries (Coarse Retrieval) You maintain a summary index, where each summary represents a larger document, chapter, or cluster.
  - Step 2 — Search Related Chunks (Fine Retrieval) Once you find the top summaries, you only search inside their associated chunks.

# Structured Corpus

- Adding meta-data: adding meta-data in the query searching to improve retrieval accuracy, provide context during chunking, and enables filtering



# Retrieval Source Optimization

---

- Adding meta-data
- Two-stage method: Retrieve documents through summaries, then retrieve text blocks from the documents.
  - Step 1 — Search Summaries (Coarse Retrieval) You maintain a summary index, where each summary represents a larger document, chapter, or cluster.
  - Step 2 — Search Related Chunks (Fine Retrieval) Once you find the top summaries, you only search inside their associated chunks.


# KG as a Retrieval Data Source

---

- Extract entities from the user's input query, then construct a subgraph to form context, and finally feed it into the large model for generation.
  - Use LLM (or other models) to extract key entities from the question.
  - Retrieve subgraphs based on entities, delving to a certain depth, such as 2 hops or even more.
  - Utilize the obtained context to generate answers through LLM.Two-stage method: Retrieve documents through summaries, then retrieve text blocks from the documents.


# KG as a Retrieval Data Source

- Extract entities from the user's input query, then construct a subgraph to form context, and finally feed it into the large model for generation.




# Query Optimization

- Questions and answers do not always possess high semantic similarity; adjusting the Query can yield better retrieval results.
- Rewrite query:



# Query Optimization


- Questions and answers do not always possess high semantic similarity; adjusting the Query can yield better retrieval results.
- Clarify the query:



# Embedding Optimization

- Better embedding always indicate a better retrieval results:
  - Selecting a more suitable embedding method
  - Fine-tuning the embedding model


Try different embedding methods in the RAG



| Model               | Retriever & Framework | Dataset     |             |             |             |             |             |             |             |
|---------------------|-----------------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
|                     |                       | HotpotQA    |             | 2Wiki       |             | NQ          |             | WebQ        |             |
|                     |                       | EM          | F1          | EM          | F1          | EM          | F1          | EM          | F1          |
| BLM25               | BLM25                 | 25.4        | 37.2        | 16.6        | 21.1        | 26.0        | 32.8        | 22.2        | 31.2        |
|                     | +SuRe                 | 38.8        | 53.5        | 23.8        | 31.0        | 36.6        | 47.9        | 34.4        | 48.5        |
|                     | +EmbQA (ours)         | <b>42.0</b> | <b>55.8</b> | <b>27.4</b> | <b>36.6</b> | <b>42.2</b> | <b>54.4</b> | <b>38.2</b> | <b>52.1</b> |
| DPR                 | DPR                   | 20.6        | 21.7        | 10.8        | 13.5        | 25.0        | 34.2        | 23.8        | 34.4        |
|                     | +SuRe                 | 25.0        | 31.9        | 14.2        | 16.0        | 38.8        | 52.3        | 36.0        | 49.6        |
|                     | +EmbQA (ours)         | <b>29.8</b> | <b>36.3</b> | <b>16.8</b> | <b>21.0</b> | <b>43.0</b> | <b>54.4</b> | <b>38.0</b> | <b>52.0</b> |
| LLaMA31.8B-Ins      | Contriever            | 22.6        | 35.4        | 16.6        | 20.7        | 25.8        | 32.8        | 25.2        | 34.2        |
|                     | +SuRe                 | 33.8        | 50.6        | 21.0        | 29.3        | 39.0        | 52.8        | 34.4        | 48.5        |
|                     | +EmbQA (ours)         | <b>36.6</b> | <b>52.7</b> | <b>26.4</b> | <b>34.2</b> | <b>42.2</b> | <b>53.6</b> | <b>36.0</b> | <b>49.6</b> |
| BLM25               | BLM25                 | 21.2        | 29.2        | 13.8        | 21.7        | 18.8        | 25.3        | 19.0        | 26.1        |
|                     | +SuRe                 | 32.2        | <b>46.1</b> | 17.8        | 30.1        | 35.2        | 45.1        | 31.6        | 45.7        |
|                     | +EmbQA (ours)         | <b>34.8</b> | <b>44.3</b> | <b>18.6</b> | <b>30.5</b> | <b>35.8</b> | <b>46.0</b> | <b>35.8</b> | <b>48.1</b> |
| DPR                 | DPR                   | 7.8         | 11.0        | 3.8         | 4.5         | 22.2        | 26.7        | 18.8        | 27.7        |
|                     | +Sure                 | 15.0        | 21.8        | 6.4         | 8.5         | 40.0        | <b>51.8</b> | 32.6        | <b>47.7</b> |
|                     | +EmbQA (ours)         | <b>16.2</b> | <b>23.3</b> | <b>7.6</b>  | <b>9.6</b>  | <b>40.2</b> | 49.4        | <b>33.4</b> | 46.0        |
| Mistral v0.2.7B-Ins | Contriever            | 19.4        | 28.6        | 13.6        | 20.7        | 21.8        | 27.4        | 17.8        | 24.4        |
|                     | +SuRe                 | 28.0        | 41.6        | 17.2        | 25.4        | 39.8        | 51.6        | 30.2        | <b>45.0</b> |
|                     | +EmbQA (ours)         | <b>29.8</b> | <b>42.3</b> | <b>17.4</b> | <b>26.2</b> | <b>40.6</b> | <b>51.8</b> | <b>31.6</b> | 43.0        |
| BLM25               | BLM25                 | 28.6        | 37.1        | 20.2        | 24.1        | 24.0        | 29.4        | 22.6        | 31.4        |
|                     | +Sure                 | 43.6        | 54.7        | 28.4        | <b>34.1</b> | 41.6        | 49.0        | 36.6        | 47.3        |
|                     | +EmbQA (ours)         | <b>44.6</b> | <b>55.6</b> | <b>28.8</b> | 33.8        | <b>42.4</b> | <b>49.2</b> | <b>38.2</b> | <b>48.7</b> |
| Owen 2.5.7B-Ins     | DPR                   | 8.8         | 9.8         | 5.6         | 7.1         | 29.2        | 32.6        | 25.6        | 31.1        |
|                     | +Sure                 | 21.8        | 27.3        | 12.2        | 16.1        | 45.4        | 54.6        | 38.4        | 49.6        |
|                     | +EmbQA (ours)         | <b>22.6</b> | <b>29.1</b> | <b>13.8</b> | <b>17.3</b> | <b>45.8</b> | <b>54.7</b> | <b>38.6</b> | <b>50.1</b> |
| Contriever          | Contriever            | 27.0        | 34.0        | 17.6        | 20.0        | 26.6        | 31.9        | 21.0        | 29.1        |
|                     | +Sure                 | 38.8        | <b>50.3</b> | 23.8        | 30.4        | 44.0        | <b>52.9</b> | 36.4        | 48.1        |
|                     | +EmbQA (ours)         | <b>39.0</b> | 50.2        | <b>24.4</b> | <b>30.9</b> | <b>45.2</b> | 50.5        | <b>37.0</b> | <b>48.6</b> |

# Embedding Optimization

- Better embedding always indicate a better retrieval results:
  - Selecting a more suitable embedding method
  - **Fine-tuning the embedding model**




## Fine-tuning on RAG

---

- It is also possible to train the RAG framework by fine-tuning, including:
  - Retriever fine-tuning
  - Generator fine-tuning

# Fine-tuning on RAG


- It is also possible to train the RAG framework by fine-tuning, including:
  - Retriever fine-tuning
  - Generator fine-tuning



Using the attention scores  
annotate which documents  
the LM “prefers”.

# Fine-tuning on RAG

- It is also possible to train the RAG framework by fine-tuning, including:
  - Retriever fine-tuning
  - Generator fine-tuning



Add an entity prediction loss in the fine-tuning

# Thank you

[Lin.1.Gui@kcl.ac.uk](mailto:Lin.1.Gui@kcl.ac.uk)

[www.kcl.ac.uk/people/lin-gui](http://www.kcl.ac.uk/people/lin-gui)